

Booklet: Introduction to Linear algebra for pytorch training

Generated by GPT-4o for Avi Salmon

October 12, 2024

1 Basics of Linear Algebra: Vectors and Matrices

Linear algebra is a fundamental area of mathematics that is essential for understanding and effectively using machine learning libraries like PyTorch. In this chapter, we will delve into the basics of linear algebra, focusing on vectors and matrices, which are the building blocks for more complex operations and algorithms.

Vectors

A vector is a mathematical object that has both magnitude and direction. In a computational context, a vector can be thought of as an array of numbers. Formally, a vector in R^n is an ordered list of numbers, written as:

$$\mathbf{v} = v_1 v_2 \dots v_n$$

Each element v_i in the vector is a component of the vector. Vectors can represent various quantities such as position, velocity, or even data features.

In PyTorch, vectors can be represented using tensors. Below is an example of how to create a simple vector:

```
import torch

# Create a 3-dimensional vector
vector = torch.tensor([1, 2, 3])
```

Operations on Vectors

Vectors support a variety of operations, including addition, scalar multiplication, and the dot product. Here, we'll briefly discuss each operation:

- **Addition:** Vectors of the same dimension can be added together by adding the corresponding components. Given two vectors **a** and **b**:

$$\mathbf{c} = a_1 + b_1 a_2 + b_2 \dots a_n + b_n$$

```
a = torch.tensor([1, 2, 3])
b = torch.tensor([4, 5, 6])
c = a + b
# c is tensor([5, 7, 9])
```

- **Scalar Multiplication:** A vector can be multiplied by a scalar (a single number), scaling each component of the vector:

$$\mathbf{b} = c \times \mathbf{a} = c \times a_1 c \times a_2 \dots c \times a_n$$

```
c = 2
b = c * a
# b is tensor([2, 4, 6])
```

- **Dot Product:** The dot product of two vectors of the same dimension is a scalar obtained by multiplying corresponding entries and summing the results:

$$\mathbf{a} \cdot \mathbf{b} = a_1 \times b_1 + a_2 \times b_2 + \cdots + a_n \times b_n$$

```
dot_product = torch.dot(a, b)
# dot_product is 32 (1*4 + 2*5 + 3*6)
```

Matrices

A matrix is a two-dimensional array of numbers arranged in rows and columns. It is often used to represent transformations and systems of linear equations. A matrix with m rows and n columns is called an $m \times n$ matrix, denoted as:

$$A = a_{11}a_{12} \cdots a_{1n} a_{21}a_{22} \cdots a_{2n} \ddots \ddots a_{m1}a_{m2} \cdots a_{mn}$$

In PyTorch, matrices are also represented using tensors, as shown below:

```
# Create a 2x3 matrix
matrix = torch.tensor([[1, 2, 3], [4, 5, 6]])
```

Operations on Matrices

Matrices support operations such as addition, scalar multiplication, and matrix multiplication:

- **Addition:** Like vectors, matrices of the same dimensions can be added together:

```
A = torch.tensor([[1, 2], [3, 4]])
B = torch.tensor([[5, 6], [7, 8]])
C = A + B
# C is tensor([[ 6,  8], [10, 12]])
```

- **Scalar Multiplication:** Each element of a matrix is multiplied by a scalar:

```
D = 2 * A
# D is tensor([[ 2,  4], [ 6,  8]])
```

- **Matrix Multiplication:** Two matrices can be multiplied if the number of columns in the first matrix equals the number of rows in the second matrix:

```
E = torch.tensor([[1, 2, 3], [4, 5, 6]])
F = torch.tensor([[7, 8], [9, 10], [11, 12]])
G = E @ F
# G is tensor([[ 58,  64], [139, 154]])
```

Practice Questions

- Create a vector with elements $[3, 5, 7]$ and perform scalar multiplication with 3. What is the resulting vector?
- Create two matrices $X = \begin{bmatrix} 12 & 34 \end{bmatrix}$ and $Y = \begin{bmatrix} 56 & 78 \end{bmatrix}$. Perform the addition and scalar multiplication of X with 5.
- Calculate the dot product of vectors $[1, 4, 6]$ and $[2, 3, 5]$.
- Multiply matrices $A = \begin{bmatrix} 24 & 13 \end{bmatrix}$ and $B = \begin{bmatrix} 37 & 58 \end{bmatrix}$.

This chapter introduces the foundational concepts of vectors and matrices, which are crucial for more advanced topics in linear algebra and machine learning. Understanding these basics will greatly simplify your journey into machine learning with PyTorch.

2 Matrix Operations and Properties

Matrix operations and properties form the backbone of many applications in linear algebra, which in turn are crucial for various computations in PyTorch. Understanding these concepts is essential for performing efficient matrix manipulations, crucial for machine learning tasks.

Matrix Addition and Subtraction

Matrix addition and subtraction are straightforward and are carried out element-wise. This means that two matrices can be added or subtracted if and only if they have the same dimensions. Consider matrices A and B of the same size:

$$A = a_{11}a_{12}a_{21}a_{22}, \quad B = b_{11}b_{12}b_{21}b_{22}$$

The addition $C = A + B$ yields:

$$C = a_{11} + b_{11}a_{12} + b_{12}a_{21} + b_{21}a_{22} + b_{22}$$

Subtraction follows similarly.

Scalar Multiplication

In scalar multiplication, each entry of the matrix is multiplied by a scalar value. Given a matrix A and a scalar c , the product cA is calculated as:

$$cA = c \cdot a_{11}c \cdot a_{12}c \cdot a_{21}c \cdot a_{22}$$

Matrix Multiplication

Matrix multiplication involves the dot product of rows of the first matrix with columns of the second matrix. Two matrices, A of size $m \times n$ and B of size $n \times p$, can be multiplied to form a matrix C of size $m \times p$.

$$C_{ij} = \sum_{k=1}^n a_{ik} \cdot b_{kj}$$

For a beginner-level example, consider:

$$A = 1234, \quad B = 5678$$

The product $C = A \cdot B$ is:

$$C = (1 \cdot 5 + 2 \cdot 7)(1 \cdot 6 + 2 \cdot 8)(3 \cdot 5 + 4 \cdot 7)(3 \cdot 6 + 4 \cdot 8) = 19224350$$

Transpose of a Matrix

The transpose of a matrix A , denoted by A^T , is obtained by swapping its rows with columns. Formally, if A is a $m \times n$ matrix, A^T is a $n \times m$ matrix. This operation is critical in linear algebra to simplify expressions and solve equations.

For example:

$$A = 1234, \quad A^T = 1324$$

Practice Questions

- Given matrices $X = \begin{bmatrix} 24 & 68 \\ 10 & 01 \end{bmatrix}$, compute $X + Y$ and $X - Y$.
- If $A = \begin{bmatrix} 35 & 79 \end{bmatrix}$, calculate $2A$ and A^T .
- Perform the matrix multiplication for $D = \begin{bmatrix} 23 & 4 \end{bmatrix}$ and $E = \begin{bmatrix} 5 \end{bmatrix}$.
- True or False: The transpose of a transpose of a matrix is the original matrix.

These foundational matrix operations and properties are integral to implementing more complex algorithms and computations within PyTorch for deep learning and data science tasks. Understanding these manipulations will enable one to model and solve real-world problems effectively.

3 Vector Spaces and Subspaces

Vectors and matrices are fundamental concepts in linear algebra, and play a crucial role in machine learning frameworks like PyTorch. In this chapter, we will delve into the concepts of vector spaces and subspaces, which form the backbone for these mathematical entities.

A **vector space** is a collection of vectors that can be added together and multiplied by scalars (numbers) to produce another vector within the same space. In a more formal sense, a vector space over a field F is a set V along with two operations: vector addition and scalar multiplication, that satisfy the following axioms (rules):

- **Closure under Addition:** For any two vectors $\mathbf{u}, \mathbf{v} \in V$, the sum $\mathbf{u} + \mathbf{v}$ is also in V .
- **Closure under Scalar Multiplication:** For any vector $\mathbf{v} \in V$ and scalar $c \in F$, the product $c\mathbf{v}$ is in V .
- **Associative and Commutative Properties:** Vector addition is both associative and commutative.
- **Additive Identity and Inverse:** There exists a zero vector $\mathbf{0}$ such that $\mathbf{v} + \mathbf{0} = \mathbf{v}$, and for every \mathbf{v} there exists a vector $-\mathbf{v}$ such that $\mathbf{v} + (-\mathbf{v}) = \mathbf{0}$.
- **Distributive Properties:** $c(\mathbf{u} + \mathbf{v}) = c\mathbf{u} + c\mathbf{v}$ and $(c + d)\mathbf{v} = c\mathbf{v} + d\mathbf{v}$ for any scalars c, d .
- **Identity Element of Scalar Multiplication:** Multiplying a vector by the scalar 1 gives the vector itself: $1\mathbf{v} = \mathbf{v}$.

Let's see this in practice using Python with PyTorch:

```
import torch

# Define two vectors in a vector space
vector_u = torch.tensor([1, 2, 3])
vector_v = torch.tensor([4, 5, 6])

# Check closure under addition
result = vector_u + vector_v
print("Addition Result: ", result)

# Check closure under scalar multiplication
scalar = 3
scaled_vector = scalar * vector_u
print("Scalar Multiplication Result: ", scaled_vector)
```

The output demonstrates the closure properties for both operations:

```
Addition Result:  tensor([5, 7, 9])
Scalar Multiplication Result:  tensor([3, 6, 9])
```

A **subspace** is simply a vector space that is contained within another vector space. In other words, it's a subset that is itself a vector space. For W to be a subspace of V , it must also satisfy the vector space axioms for addition and scalar multiplication, and contain the zero vector of V .

To check if a set is a subspace, you can use these three criteria:

- The zero vector of V is in W .
- W is closed under vector addition.
- W is closed under scalar multiplication.

Consider the example using PyTorch:

```
# Subspace check examples

# Define some vectors
vector_zero = torch.tensor([0, 0, 0])
vector_w1 = torch.tensor([1, 2, 3])
vector_w2 = torch.tensor([-1, -2, -3])

# Sum of vectors in the potential subspace
sum_vectors = vector_w1 + vector_w2

# Scalar multiplication of a vector in the potential subspace
scalar_multiplication = 2 * vector_w1

print("Sum of vectors:", sum_vectors)
print("Scalar Multiplication:", scalar_multiplication)
```

Expected output:

```
Sum of vectors: tensor([0, 0, 0])
Scalar Multiplication: tensor([2, 4, 6])
```

In these examples, the sum of \mathbf{w}_1 and \mathbf{w}_2 results in the zero vector, and scalar multiplication results in another vector in the potential subspace, indicating closure under the operations.

Practice Questions

- Given the set of all vectors x and y where $x = 2y$, determine whether this set forms a subspace of R^2 .
- Consider the vectors $\mathbf{u} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$ and $\mathbf{v} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ in R^2 . Show that the span of these vectors form the complete vector space R^2 .
- Verify if the set of all matrices of the form $\begin{bmatrix} a & b \\ 0 & c \end{bmatrix}$ is a subspace of the space of all 2×2 matrices.

4 Linear Transformations

Linear transformations are fundamental concepts in linear algebra that enable us to describe how vectors and spaces are manipulated through linear mappings. In simple terms, a linear transformation is a function between vector spaces that preserves the operations of vector addition and scalar multiplication.

4.1 Understanding Linear Transformations

A linear transformation T from a vector space V to a vector space W is a rule that assigns to each vector \mathbf{v} in V a unique vector $T(\mathbf{v})$ in W . The key properties that characterize a linear transformation are:

- **Additivity:** $T(\mathbf{u} + \mathbf{v}) = T(\mathbf{u}) + T(\mathbf{v})$ for all vectors $\mathbf{u}, \mathbf{v} \in V$.
- **Homogeneity of Degree 1:** $T(c\mathbf{v}) = cT(\mathbf{v})$ for all vectors $\mathbf{v} \in V$ and scalars c .

These properties ensure that the transformation respects the structure of the vector space.

4.2 Matrix Representation of Linear Transformations

In practice, linear transformations are often represented using matrices. If $T : R^n \rightarrow R^m$ is a linear transformation, there exists a unique $m \times n$ matrix A such that for every vector $\mathbf{x} \in R^n$, the transformation is given by matrix multiplication:

$$T(\mathbf{x}) = A\mathbf{x}$$

Here, the matrix A acts as a transformation operator on the vector \mathbf{x} , transforming it into another vector in R^m .

4.3 Example of a Linear Transformation

Consider a linear transformation $T : R^2 \rightarrow R^2$ defined by the matrix

$$A = \begin{bmatrix} 2 & 3 \\ 1 & 4 \end{bmatrix}$$

For a vector $\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$, the transformation is given by:

$$T(\mathbf{x}) = A\mathbf{x} = \begin{bmatrix} 2 & 3 \\ 1 & 4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 2x_1 + 3x_2 \\ x_1 + 4x_2 \end{bmatrix}$$

This results in a new vector formed by linearly transforming \mathbf{x} .

4.4 Implementing Linear Transformations in Python with PyTorch

With PyTorch, linear transformations can be easily implemented using the ‘torch’ library. Below is an example illustrating how to perform a linear transformation in PyTorch:

```
import torch

# Define the transformation matrix
A = torch.tensor([[2, 3],
                  [1, 4]])

# Define a vector
x = torch.tensor([[1], [2]])

# Apply the linear transformation
y = torch.matmul(A, x)

print(y)
```

In this code snippet, the matrix A and vector x are defined using PyTorch tensors. The function ‘torch.matmul’ is used to perform matrix multiplication, applying the linear transformation.

4.5 Practice Questions

Here are some practice questions to solidify your understanding of linear transformations:

- Consider the linear transformation $T : R^3 \rightarrow R^2$ given by a matrix $B = \begin{bmatrix} 10 & 2 \\ 13 & 0 \end{bmatrix}$. Compute $T(\mathbf{v})$ for $\mathbf{v} = \begin{bmatrix} 2 \\ -1 \\ 3 \end{bmatrix}$.
- Given a linear transformation $T : R^2 \rightarrow R^2$ represented by the matrix $C = \begin{bmatrix} 0 & 1 \\ -10 & 0 \end{bmatrix}$, describe the effect of T on vectors in R^2 .
- Implement a linear transformation in PyTorch where the transformation matrix is $D = \begin{bmatrix} 3 & -1 \\ 20 & 0 \end{bmatrix}$ and apply it to the vector $\mathbf{w} = \begin{bmatrix} 4 \\ 1 \end{bmatrix}$. Print the resulting vector.

Understanding linear transformations and their representations through matrices is crucial for working with linear algebra in machine learning, specifically when dealing with neural networks and other data transformations. Practice these concepts to deepen your comprehension and application in PyTorch.

5 System of Linear Equations

A system of linear equations consists of two or more linear equations involving the same set of variables. The primary objective when dealing with these systems is to find the values of the variables that simultaneously satisfy all the equations in the system.

Introduction to Systems of Linear Equations

Linear equations represent lines when plotted on a graph. When dealing with two variables, these can be expressed in the general form:

$$ax + by = c$$

where a , b , and c are constants. A system of linear equations in two variables might look like:

$$a_1x + b_1y = c_1 \quad a_2x + b_2y = c_2$$

The solution to this system is the point (x, y) where the two lines intersect on a graph.

Solving Systems of Linear Equations

There are several methods to solve a system of linear equations:

- **Graphical Method:** Plot each equation on a graph and identify the point(s) of intersection.
- **Substitution Method:** Solve one of the equations for one variable and substitute that expression in the other equation.
- **Elimination Method:** Add or subtract equations to eliminate one variable, allowing you to solve for the remaining variable.
- **Matrix Method:** Use matrix operations (as discussed in earlier chapters) to solve the system. This is particularly useful for larger systems.

Example: Solving with the Substitution Method

Consider the system:

$$3x + 4y = 14 \quad 2x - y = 1$$

Solve the second equation for y :

$$y = 2x - 1$$

Substitute y in the first equation:

$$3x + 4(2x - 1) = 14$$

Simplify and solve for x :

$$3x + 8x - 4 = 1411x = 18x = \frac{18}{11}$$

Now find y using $x = \frac{18}{11}$:

$$y = 2 \left(\frac{18}{11} \right) - 1 = \frac{36}{11} - \frac{11}{11} = \frac{25}{11}$$

Thus, the solution is:

$$(x, y) = \left(\frac{18}{11}, \frac{25}{11} \right)$$

Using Technology to Solve Systems

Python, and particularly libraries such as NumPy, make solving systems of equations more streamlined. Here's how you can solve the system using Python in cmd line using the Windows command prompt to run a Python script:
in cmd line:

```
> python -c "import numpy as np; a = np.array([[3, 4], [2, -1]]); b
          = np.array([14, 1]);
> x = np.linalg.solve(a, b); print('Solution:', x)"
```

This code uses NumPy's `linalg.solve` function to solve the system, yielding the solution: (x, y) .

Practice Questions

- Solve the following system of equations using the substitution method:

$$x + 3y = 92x - y = 4$$

- Solve the system using the elimination method:

$$4x - 2y = 6 - 2x + 3y = -3$$

- Using Python and the method shown above, solve:

$$5x + 2y = 18 - 3x + 4y = 3$$

Interpret the solution and graph the equations to verify visually.

6 Eigenvalues and Eigenvectors

Understanding eigenvalues and eigenvectors is crucial for many applications in linear algebra, especially when working with machine learning libraries like PyTorch. In this chapter, we will introduce the concepts of eigenvalues and eigenvectors, explain how to compute them, and illustrate their importance with code examples.

Eigenvalues and Eigenvectors: Introduction

In linear algebra, eigenvalues and eigenvectors are properties of a square matrix. They offer insights into the matrix's characteristics and are used in various computations, such as matrix decomposition, which is vital in machine learning algorithms.

Matrices can be thought of as transformations applied to vectors. An eigenvector of a given square matrix is a vector whose direction remains unchanged when the matrix is applied to it. The scalar multiple by which the eigenvector is stretched or compressed is called the eigenvalue.

Mathematically, for a square matrix A , a non-zero vector v is an eigenvector if:

$$A \cdot v = \lambda \cdot v$$

where λ is the eigenvalue associated with the eigenvector v .

Computing Eigenvalues and Eigenvectors

To compute eigenvalues and eigenvectors, we typically solve the characteristic equation:

$$\det(A - \lambda I) = 0$$

where I is the identity matrix of the same size as A . Solving this equation gives us the eigenvalues, and substituting each eigenvalue back into the equation $A \cdot v = \lambda \cdot v$ allows us to find the corresponding eigenvectors.

Though manual computation can be challenging for large matrices, libraries like NumPy and PyTorch simplify this process.

Example: Computing Eigenvalues and Eigenvectors Using PyTorch

The following example shows how to compute eigenvalues and eigenvectors in PyTorch. First, ensure that you have PyTorch installed on your Windows system.

To install PyTorch, in cmd line:

```
> pip install torch
```

Once installed, you can compute eigenvalues and eigenvectors as follows:

```
import torch

# Define a square matrix
A = torch.tensor([[2.0, 1.0],
                  [1.0, 2.0]])

# Compute eigenvalues and eigenvectors
eigenvalues, eigenvectors = torch.eig(A, eigenvectors=True)
```

```
print("Eigenvalues:")
print(eigenvalues)
print("Eigenvectors:")
print(eigenvectors)
```

This script defines a 2x2 matrix A , computes its eigenvalues and eigenvectors using the `torch.eig()` function, and prints the results.

Applications of Eigenvalues and Eigenvectors

Eigenvalues and eigenvectors are used to understand the geometric properties of transformations represented by matrices. They are crucial in:

- Principal Component Analysis (PCA) for dimensionality reduction.
- Optimization problems in machine learning.
- Stability analysis in differential equations and dynamic systems.

Practice Questions

- Find the eigenvalues and eigenvectors of the matrix $\begin{bmatrix} 4 & 1 \\ 2 & 3 \end{bmatrix}$.
- Explain why eigenvectors are critical in understanding matrix operations.
- Using PyTorch, solve for the eigenvalues and eigenvectors of a 3x3 matrix of your choice. Discuss your results.

Eigenvalues and eigenvectors provide significant insights into the behavior of linear transformations and are vital tools for advanced topics in machine learning and other scientific computations. Understanding these concepts bridges linear algebra and practical applications in data science.

7. Singular Value Decomposition (SVD)

Singular Value Decomposition (SVD) is a fundamental matrix factorization technique in linear algebra. It is used widely in various applications, including machine learning, image processing, and data compression. This chapter will introduce the concept of SVD, demonstrate its significance, and provide practical examples for beginners using PyTorch.

Understanding Singular Value Decomposition

SVD decomposes any given $m \times n$ matrix A into three other matrices: U , Σ , and V^T . Mathematically, this is represented as:

$$A = U\Sigma V^T$$

Here,

- U is an $m \times m$ orthogonal matrix.
- Σ is an $m \times n$ diagonal matrix with non-negative real numbers on the diagonal.
- V^T is an $n \times n$ orthogonal matrix.

The columns of U are called the left singular vectors, the columns of V are called the right singular vectors, and the diagonal values of Σ are the singular values.

Significance of SVD

SVD is a powerful tool due to the insightful information it provides about the original matrix. Some notable applications include:

- **Dimensionality Reduction:** By keeping only a few significant singular values, we can reduce the dimensionality of data while preserving most of its information. This is often used in Principal Component Analysis (PCA).
- **Data Compression:** Using SVD, data can be compressed by representing it with lower rank approximations.
- **Noise Reduction:** In image processing, SVD can be used to separate noise from underlying data structures.

Computing SVD in PyTorch

PyTorch provides a simple function to compute the SVD of a matrix. Below is an example of how to use PyTorch to compute the SVD of a matrix:

```
import torch

# Create a random matrix
A = torch.tensor([[3.0, 2.0, 2.0], [2.0, 3.0, -2.0]])

# Compute SVD
```

```

U, S, V = torch.svd(A)

print("U:\n", U)
print("Singular values:\n", S)
print("V:\n", V)

```

Listing 1: Python code to compute SVD using PyTorch

In this code snippet, we first import the PyTorch library and create a random matrix A . We then use the `torch.svd()` function to obtain matrices U , S , and V .

Practice Questions

- Compute the SVD of the following matrix using PyTorch:

$$403 - 5$$

Use the results to reconstruct the original matrix.

- Explain how SVD can be used to improve the performance of machine learning models. Give examples of scenarios where SVD might be beneficial.
- Given a matrix:

$$A = 122334$$

Calculate its SVD using PyTorch and plot the singular values. Discuss any observation you make from the singular value plot.

- Investigate how to approximate a matrix by using fewer singular values. Compute the SVD of a matrix of your choice and test approximating it by using only the top 1 or 2 singular values. Discuss the impact on matrix accuracy and size.

This chapter introduces the basics of Singular Value Decomposition (SVD) and its implementation using PyTorch, providing a foundation for exploring more advanced applications and techniques in linear algebra numerical computations.

8 Introduction to PyTorch: Tensors

PyTorch is a popular open-source machine learning library that provides a flexible and efficient platform for deep learning research and production. At its core, PyTorch leverages a key data structure called **Tensors**, which are similar to NumPy arrays but with additional features that support GPU acceleration, automatic differentiation, and more. This chapter provides an introduction to Tensors in PyTorch, highlighting their operations and usage.

What are Tensors?

Tensors can be thought of as generalized matrices. In mathematics, a tensor extends the idea of scalars (0D), vectors (1D), and matrices (2D) to higher dimensions. In PyTorch, Tensors are used to store data that can be manipulated in a variety of ways.

Creating Tensors

You can create Tensors in several ways:

- Directly from data.
- From NumPy arrays.
- Using PyTorch’s built-in functions.

Here is how you create a Tensor using PyTorch:

```
import torch

# Creating a tensor from a list
data = [[1, 2], [3, 4]]
tensor_from_list = torch.tensor(data)

# Creating a tensor from a NumPy array
import numpy as np
numpy_array = np.array(data)
tensor_from_numpy = torch.from_numpy(numpy_array)

# Creating a tensor using built-in functions
tensor_zeros = torch.zeros((2, 2))
tensor_ones = torch.ones((2, 2))
tensor_random = torch.rand((2, 2))
```

Tensor Operations

Tensors support a wide range of operations, just like NumPy arrays. You can perform arithmetic operations, slice tensors, and more.

```
# Arithmetic operations
a = torch.tensor([2, 3])
b = torch.tensor([5, 6])
```

```

# Addition
sum_tensor = a + b

# Element-wise multiplication
product_tensor = a * b

# Dot product
dot_product = torch.dot(a, b)

# Slicing tensors
matrix = torch.tensor([[1, 2, 3], [4, 5, 6]])
# Extract the first row
first_row = matrix[0, :]
# Extract the second column
second_column = matrix[:, 1]

```

Tensor Properties

Tensors have several important properties:

- **Shape:** The dimensions of the tensor.
- **Data type:** The data type of the tensor elements, such as `float32`, `int64`, etc.
- **Device:** Tensors can reside on CPU or GPU, which is essential for leveraging hardware acceleration.

```

# Checking tensor properties
tensor = torch.rand((3, 3))

# Shape of the tensor
print(tensor.shape)

# Data type of the tensor
print(tensor.dtype)

# Device of the tensor
print(tensor.device)

```

Moving Tensors to GPU

To take advantage of GPU acceleration, you can move Tensors to GPU. Note that PyTorch must be installed with CUDA support for this.

In cmd line:

```
>pip install torch torchvision torchaudio --extra-index-url https://download.pytorch.org
```

```

# Check if GPU is available and move tensor to GPU
device = torch.device('cuda' if torch.cuda.is_available() else
                      'cpu')

```

```
tensor = tensor.to(device)
```

Practice Questions

Here are a few practice questions to deepen your understanding of Tensors and their operations:

- Create a 3x3 Tensor filled with random numbers and compute its transpose.
- Given two random Tensors, compute their mean and standard deviation.
- Move a Tensor from CPU to GPU, perform matrix multiplication, and move it back to CPU. Verify the result by comparing with a CPU-only computation.
- Using slicing, extract a sub-tensor from a given larger Tensor and perform an element-wise operation on it.

Understanding and manipulating Tensors efficiently is crucial for leveraging PyTorch in machine learning tasks. This foundation will help you undertake more advanced topics in PyTorch and deep learning.

9 - PyTorch Tensor Operations

When working with PyTorch, tensors are the central data structure. Tensors are similar to NumPy arrays, but they come with additional functionality for constructing and training deep learning models. This chapter focuses on performing operations with these tensors, equipping you with the skills to manipulate and transform data effectively in PyTorch.

Basic Tensor Operations

Once you've grasped the concept of tensors, you'll want to perform various operations on them. This includes arithmetic operations, reshaping, and transposing, among others. Let's start by looking at some basic tensor operations in PyTorch.

In order to proceed, ensure that PyTorch is installed on your Windows machine. You can install it using PyPI (Python Package Index) as follows: in cmd line:

```
> pip install torch
```

Below is a Python code snippet demonstrating basic arithmetic operations with tensors:

```
import torch

# Create two tensors
a = torch.tensor([1, 2, 3])
b = torch.tensor([4, 5, 6])

# Basic arithmetic
sum_tensor = a + b
diff_tensor = a - b
prod_tensor = a * b
quo_tensor = a / b

print("Sum: ", sum_tensor)
print("Difference: ", diff_tensor)
print("Product: ", prod_tensor)
print("Quotient: ", quo_tensor)
```

In this example, we used basic arithmetic operators such as addition, subtraction, multiplication, and division directly on tensors. Each operation is element-wise, meaning it is performed element by element across corresponding positions in the tensors.

Tensor Reshaping and Manipulation

Sometimes you may need to change the shape of a tensor without altering its data. One common operation is reshaping, which can be performed using the `.view()` or `.reshape()` method. The latter is more versatile

in recent PyTorch versions, automatically calculating dimensions when possible.

Here's how you can reshape a tensor:

```
# Create a tensor with 8 elements
original_tensor = torch.tensor([1, 2, 3, 4, 5, 6, 7, 8])

# Reshape to a 2x4 matrix
reshaped_tensor = original_tensor.reshape(2, 4)
print("Reshaped Tensor:\n", reshaped_tensor)
```

In the code above, a 1D tensor with 8 elements is reshaped into a 2D tensor with dimensions 2x4.

Tensor Transposition

Transposing is another manipulation technique where the dimensions of a tensor are swapped. In a 2D tensor (matrix), this operation switches rows and columns. You can transpose a tensor using the `.t()` method for 2D tensors or `.transpose()` for higher dimensional tensors.

```
# Create a 2x3 tensor
matrix = torch.tensor([[1, 2, 3], [4, 5, 6]])

# Transpose the matrix
transposed_matrix = matrix.t()
print("Transposed Matrix:\n", transposed_matrix)
```

Other Useful Operations

There are numerous other tensor operations, such as finding the maximum or minimum value, computing the mean, and summation. Here are a few examples:

```
# Create a 1D tensor
tensor = torch.tensor([5, 3, 8, 1])

# Find the maximum value and its index
max_val, max_idx = tensor.max(0)
print("Maximum Value: ", max_val, " at Index: ", max_idx.item())

# Compute the mean
mean_val = tensor.float().mean()
print("Mean Value: ", mean_val)

# Sum of the tensor
sum_val = tensor.sum()
print("Sum: ", sum_val)
```

These operations are particularly useful when processing batches of data in machine learning. Understanding how to effectively manipulate tensors will aid significantly in training models and handling large datasets.

Practice Questions

- Create two PyTorch tensors with arbitrary integer values. Perform element-wise addition, subtraction, multiplication, and division on them. Print the results.
- Given a 1D tensor with 12 elements, reshape it into a 3D tensor with dimensions $3 \times 2 \times 2$. Print the reshaped tensor.
- Create a 4×4 tensor with random integers between 0 and 10. Transpose the tensor and print the result. What do you observe?
- Calculate the maximum, minimum, and mean values of a random tensor of size 10. Also, print the indices of the maximum and minimum values.

These exercises will help reinforce your understanding of tensor operations in PyTorch. Experimenting with different tensor manipulations forms a strong foundation for building and training neural networks.

10 Implementing Matrix Multiplication in PyTorch

Introduction

Matrix multiplication is a fundamental operation in linear algebra, commonly used in various applications including machine learning, computer graphics, and scientific computing. In PyTorch, which is a popular deep learning library, matrix multiplication can be performed efficiently, leveraging both CPU and GPU resources for enhanced performance. This chapter will guide you through the basics of matrix multiplication in PyTorch, focusing on implementation details suited for beginners.

Understanding Matrix Multiplication

Matrix multiplication involves two matrices, say **A** and **B**, where the number of columns in **A** must equal the number of rows in **B**. The result is a new matrix where each element is obtained by computing the dot product of corresponding row vectors from **A** with column vectors from **B**.

Consider two matrices:

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{21} & a_{22} \end{bmatrix}, \quad B = \begin{bmatrix} b_{11} & b_{12} & b_{21} & b_{22} \end{bmatrix}$$

The matrix product $C = A * B$ is given by:

$$C = \begin{bmatrix} a_{11}b_{11} + a_{12}b_{21} & a_{11}b_{12} + a_{12}b_{22} & a_{21}b_{11} + a_{22}b_{21} & a_{21}b_{12} + a_{22}b_{22} \end{bmatrix}$$

Matrix Multiplication in PyTorch

PyTorch provides a simple yet powerful way to perform matrix multiplication using the `torch.matmul` function or the `@` operator for tensors. Below is a step-by-step guide and example code demonstrating matrix multiplication in PyTorch.

Example Code

First, install PyTorch on your Windows system. Make sure Python is installed, and then in cmd line:

```
> pip install torch
```

Here is a code example that demonstrates how to perform matrix multiplication using PyTorch:

```
# Import the torch library
import torch

# Define matrix A (2x3)
A = torch.tensor([[1, 2, 3],
                  [4, 5, 6]])
```

```

# Define matrix B (3x2)
B = torch.tensor([[7, 8],
                  [9, 10],
                  [11, 12]])

# Perform matrix multiplication
C = torch.matmul(A, B)

print("Matrix A:")
print(A)

print("Matrix B:")
print(B)

print("Result of A * B:")
print(C)

```

Listing 2: Matrix Multiplication using PyTorch

Explanation

- We begin by importing the PyTorch library. - Matrix A is initialized as a 2x3 tensor. - Matrix B is initialized as a 3x2 tensor. - The `torch.matmul()` function is used to multiply A and B, yielding a 2x2 matrix. - The result is printed to the console.

Alternative Methods

You can also use the `@` operator for matrix multiplication, as shown below:

```

# Perform matrix multiplication using the @ operator
C = A @ B

print("Result of A @ B:")
print(C)

```

Listing 3: Using `@` operator for Matrix Multiplication

Both methods are interchangeable, and using `@` is often more concise.

Practice Questions

- Create two matrices X and Y with dimensions 4x2 and 2x3, respectively. Perform the matrix multiplication and verify the dimensions of the result.
- Using PyTorch, multiply two identity matrices of size 3x3. Verify that the result is also an identity matrix.
- Modify the matrix multiplication example to use tensors of shape 3x3 and calculate the product in PyTorch. Explain the result.

This chapter introduced you to the basics of implementing matrix multiplication in PyTorch. Mastering these fundamentals opens up a wide range of applications, especially in deep learning contexts where matrix operations are pivotal.

11 Using PyTorch for Solving Linear Systems

In this chapter, we will explore how to utilize PyTorch for solving linear systems. This involves setting up systems of linear equations, using PyTorch's functionalities to find solutions, and understanding these solutions in the context of linear algebra. PyTorch provides efficient tools for this purpose, making it an excellent choice for machine learning and scientific computing tasks that require solving such systems.

Introduction to Linear Systems

A linear system is a collection of linear equations that we solve simultaneously. In matrix form, a linear system can be represented as:

$$Ax = b$$

Where A is a matrix of coefficients, x is the vector of unknowns, and b is the result vector. The goal is to find the vector x .

Solving Linear Systems in PyTorch

PyTorch provides the `torch.linalg.solve()` function, which computes the solution of a square system of linear equations with a given matrix A and a result vector b .

Example: Solving a Linear System in PyTorch

Consider solving the following system of equations:

$$2x + 3y = 54x + 6y = 10$$

This system is represented in matrix form as:

$$A = \begin{bmatrix} 2 & 3 \\ 54 & 6 \end{bmatrix}, \quad b = \begin{bmatrix} 10 \\ 10 \end{bmatrix}$$

Here's how you can solve the linear system using PyTorch:

```
import torch

# Define the matrix A and vector b
A = torch.tensor([[2.0, 3.0], [54.0, 6.0]])
b = torch.tensor([10.0, 10.0])

# Solve the linear system Ax = b
x = torch.linalg.solve(A, b)

print(x)
```

This code will print the solution vector x . It's important to note that this particular system has infinitely many solutions since the matrix A is singular (its rows are linearly dependent).

Checking the Solution

To verify the solution obtained, you can multiply the matrix A by the solution vector x and check if it equals the original vector b .

```
# Verify the solution
b_computed = torch.matmul(A, x)
print(b_computed)
```

The output should be approximately equal to the original vector b , validating the computed solution.

Practice Questions

- Solve the following system of equations using PyTorch:

$$3x + 2y - z = 12x - 2y + 4z = -2 - x + 0.5y - z = 0$$

- Verify if the matrix $A = \begin{bmatrix} 12 & -2 & 4 \\ -2 & -1 & 0.5 \\ -1 & 0.5 & -1 \end{bmatrix}$ is singular using PyTorch, and discuss the implications for solving the system $Ax = b$.
- Consider a system of equations with no solutions. What would be the PyTorch response when attempting to solve such a system? Provide an example by creating and solving the system.

These exercises will help deepen your understanding of how to solve linear systems using PyTorch, and prepare you for more complex applications in data science and machine learning.

12 PyTorch Autograd and Backpropagation

Understanding PyTorch’s autograd and backpropagation is crucial for efficiently designing and training models in machine learning. This chapter will introduce you to these essential concepts with straightforward explanations and examples.

Introduction to Autograd

PyTorch provides an automatic differentiation system called **autograd**. It tracks all operations on `torch.Tensor` objects with the `requires_grad` attribute set to `True`. By doing so, it creates a computation graph, which enables efficient computation of gradients—needed for optimizing models via backpropagation. This makes PyTorch a preferred choice for tasks involving deep learning.

Computational Graphs

In PyTorch, each tensor operation creates a node in a computational graph. Tensors with `requires_grad=True` will be traced, and the entire history of computation will be captured in the graph. This graph is dynamic and automatically adjusted as tensors become unused (e.g., they go out of scope), which can be more efficient than static graph systems.

Example: Simple Autograd Usage

Let’s see a basic example of using autograd in PyTorch:

```
import torch

# Create tensors
x = torch.randn(3, requires_grad=True)
y = x * 2
z = y.mean()

# Backpropagation
z.backward()

# Gradients
print(x.grad)
```

In this example, we:

- Created a tensor `x` with `requires_grad=True`.
- Performed operations on `x` to build a computational graph.
- Used `backward()` on `z` to compute the gradient of `z` with respect to `x`.
- Accessed `x.grad` to view the computed gradients.

Backpropagation Explained

Backpropagation is an algorithm used to compute the gradient of a loss function with respect to all weights in the network. It works by chain rule calculation across layers in the network, allowing it to efficiently compute gradients for optimization algorithms like gradient descent.

In PyTorch, once you call `.backward()` on the final output tensor, PyTorch automatically computes the gradient of all tensors involved in graph.

Practice Questions

- Create a tensor with shape `(2, 2)` with `requires_grad=True`. Perform a series of operations to calculate a scalar output and compute the gradient. What is the resulting gradient?
- Use the concept of autograd to compute the gradient descent optimization for a function $f(x) = (x - 3)^2$. *Use PyTorch to implement a simple loop to minimize $f(x)$ starting from $x = 5$.*
- Explain the role of the attribute `.grad` in PyTorch, and how it relates to tensors used in neural networks.

By understanding these foundational concepts of autograd and backpropagation in PyTorch, you will be better equipped to design and train neural networks efficiently. As you proceed with further chapters, you will frequently encounter these principles in more complex scenarios.

13 Understanding Neural Networks as Linear Transformations

Understanding Linear Transformations in Neural Networks

Neural networks, at their core, can be viewed as a series of linear transformations followed by non-linear activations. Understanding this fundamental principle allows us to use linear algebra and PyTorch to develop and train efficient models. Linear transformations in neural networks are often represented by matrix multiplications.

Suppose we have an input vector \mathbf{x} and a weight matrix \mathbf{W} . The product of \mathbf{W} and \mathbf{x} results in a new vector \mathbf{y} , which can be expressed as:

$$\mathbf{y} = \mathbf{W} \times \mathbf{x}$$

This transformation changes the space and dimensions of the input data, which is a significant aspect of how neural networks operate.

Matrix Multiplication in Neural Networks

In the context of a simple neural network layer, the operations performed can be broken down into:

- \mathbf{W} is the weight matrix associated with the layer. - \mathbf{b} is the bias vector. - \mathbf{x} is the input vector.

The output \mathbf{y} is computed by the expression:

$$\mathbf{y} = \mathbf{W} \times \mathbf{x} + \mathbf{b}$$

This operation combines linear transformation with an addition of the bias term. The application of a non-linear activation function, such as ReLU (Rectified Linear Unit), transforms the output further to introduce non-linearity:

$$\mathbf{a} = \text{ReLU}(\mathbf{y})$$

Below is an example demonstrating this concept using PyTorch:

```
import torch
import torch.nn as nn

# Inputs and weights
x = torch.tensor([1.0, 2.0])
W = torch.tensor([[0.5, -1.0], [1.0, 0.5]])
b = torch.tensor([0.0, 1.0])

# Linear transformation
y = torch.matmul(W, x) + b
```

```
# Activation function (ReLU)
a = nn.ReLU()(y)

print("Output after linear transformation:", y)
print("Output after ReLU activation:", a)
```

Neural Network as a Series of Linear Transformations

A neural network is composed of multiple layers, each performing its linear transformation. These layers, when stacked together, enable the model to learn complex patterns.

Consider a two-layer neural network. The first layer transforms the input, and the second layer transforms the output of the first layer. Mathematically, this can be expressed as:

$$\mathbf{h} = \mathbf{W}_1 \times \mathbf{x} + \mathbf{b}_1$$

$$\mathbf{a} = \text{ReLU}(\mathbf{h})$$

$$\mathbf{o} = \mathbf{W}_2 \times \mathbf{a} + \mathbf{b}_2$$

The combined transformation allows for a more expressive model capable of solving complex problems.

Practice Questions

- Experiment with changing the values of the weight matrix \mathbf{W} and bias vector \mathbf{b} in the provided code example. Observe how the output changes. What impact does each parameter have on the final result?
- Implement a simple two-layer neural network using PyTorch, with custom weight and bias values for each layer. Apply a non-linear activation function between the layers and compute the final output.
- Analyze the effect of different activation functions (besides ReLU), such as sigmoid or tanh, on the transformed output in a neural network layer.
- Create a one-layer linear transformation model without an activation function. Train it on a simple dataset to observe the model's capability to linearly separate data points.

Understanding neural networks through the lens of linear transformations provides a clear insight into how they process and learn from data. This perspective is foundational for further in-depth study of neural network architectures and optimization techniques.“ latex

14 Dimensionality Reduction with PyTorch

Dimensionality reduction is a critical technique in data science and machine learning, allowing us to simplify datasets and extract valuable insights by reducing the number of variables under consideration. In this chapter, we will explore how to implement dimensionality reduction using PyTorch, focusing on intuitive concepts and practical implementations.

Understanding Dimensionality Reduction

Dimensionality reduction is the process of reducing the number of random variables under consideration by obtaining a set of principal variables. It is typically classified into two categories:

- **Feature Selection:** Involves selecting a subset of the original variables.
- **Feature Extraction:** Involves creating new variables by combining the original variables to form a reduced dataset.

One of the primary motivations for dimensionality reduction is to tackle the "curse of dimensionality," which suggests that high-dimensional datasets are often difficult to analyze and visualize.

Principal Component Analysis (PCA) in PyTorch

Principal Component Analysis (PCA) is one of the most popular techniques for dimensionality reduction, especially in feature extraction. It transforms the data into a new coordinate system, reducing the number of dimensions while preserving as much variance as possible.

In PyTorch, implementing PCA can be simplified as follows:

```
import torch

# Create a tensor representing a dataset with high
# dimensionality
data = torch.tensor([[2.5, 2.4],
                     [0.5, 0.7],
                     [2.2, 2.9],
                     [1.9, 2.2],
                     [3.1, 3.0],
                     [2.3, 2.7],
                     [2, 1.6]], dtype=torch.float32)

# Calculate the mean of the dataset
mean = torch.mean(data, axis=0)

# Center the dataset
centered_data = data - mean

# Compute the covariance matrix
covariance_matrix = torch.mm(centered_data.t(), centered_data)
                    / (centered_data.size(0) - 1)
```

```
# Perform Singular Value Decomposition
U, S, V = torch.svd(covariance_matrix)

# Project the data onto the new basis
pca_data = torch.mm(centered_data, U[:, :1])
print(pca_data)
```

This code snippet demonstrates how to perform PCA on a dataset using PyTorch's tensor operations. The dataset is first centered by subtracting the mean, and the covariance matrix is computed. Singular Value Decomposition (SVD) is then applied to find the principal components.

Applications of Dimensionality Reduction

Dimensionality reduction can be applied to:

- **Data Visualization:** Simplifying data to enable easy visualization in 2D or 3D.
- **Noise Reduction:** Removing noise to improve model accuracy.
- **Feature Engineering:** Creating targeted features that help improve model predictions.

By reducing the dimensions of the data, we can help prevent overfitting in machine learning models and reduce computational costs.

Practice Questions

- What are the differences between feature selection and feature extraction in dimensionality reduction?
- Implement a simple example of PCA using another dataset in PyTorch.
- How can dimensionality reduction be useful in deep learning and neural networks?
- Explain how the curse of dimensionality can affect machine learning models and how dimensionality reduction can alleviate these issues.
- Try performing PCA on an image dataset using PyTorch, and visualize the results.

““

15 Practical Applications of Linear Algebra in PyTorch Models

Linear algebra serves as the mathematical foundation for many machine learning algorithms and models, including those implemented in PyTorch. Understanding how linear algebra concepts are applied in PyTorch models can enhance your ability to build and optimize machine learning systems effectively.

Applications in Neural Networks

Neural networks, the backbone of many modern machine learning applications, inherently rely on linear algebra. At a high level, a neural network can be thought of as a series of linear transformations followed by non-linear activation functions. Each layer in a neural network applies a linear transformation to its input, combining it with the weights of the network:

$$y = Wx + b$$

where W is the weight matrix, x is the input tensor, b is the bias vector, and y is the output tensor after the linear transformation.

Example: Implementing a Simple Neural Layer in PyTorch

Let's see how we can implement a single layer of a neural network in PyTorch. This example demonstrates a simple linear transformation followed by a non-linear activation function, using the ReLU (Rectified Linear Unit) function:

```
import torch
import torch.nn.functional as F

# Define the size of input and output
input_size = 4
output_size = 3

# Create random input tensor
x = torch.randn(1, input_size)

# Initialize weights and bias
W = torch.randn(input_size, output_size)
b = torch.randn(output_size)

# Perform linear transformation
y = torch.matmul(x, W) + b

# Apply ReLU activation function
output = F.relu(y)

print(output)
```

In this code, `torch.matmul()` performs the matrix multiplication of the input tensor `x` and the weight matrix `W`. After adding the bias `b`, the ReLU activation function is applied to introduce non-linearity.

Dimensionality Reduction

Dimensionality reduction is another area where linear algebra plays a critical role, particularly through techniques like Principal Component Analysis (PCA). In PyTorch, you can perform dimensionality reduction using SVD (Singular Value Decomposition) and other methods to simplify the dataset while retaining its essential features.

Example: Dimensionality Reduction using SVD in PyTorch

The following example demonstrates how to perform SVD in PyTorch to reduce the dimensions of a dataset:

```
import torch
import torch.linalg as linalg

# Create a random matrix
data_matrix = torch.randn(100, 50)

# Perform Singular Value Decomposition
U, S, V = linalg.svd(data_matrix)

# Reduce dimensions
k = 10 # number of dimensions to keep
reduced_data = torch.matmul(U[:, :k], torch.diag(S[:k]))

print(reduced_data.shape)
```

This code snippet extracts the top k singular values and corresponding vectors, reducing the dataset to a smaller dimension while preserving its structure.

Practice Questions

- Explain how matrix multiplication is applied in the forward pass of a neural network.
- Write a PyTorch function to perform a linear transformation on a given input tensor without using `torch.matmul`. Use basic tensor operations.
- Implement a custom PyTorch function that performs ReLU activation on an input tensor without using `torch.nn.functional.relu`.
- Describe a practical scenario where dimensionality reduction could be beneficial in a machine learning project.

These exercises will help solidify your understanding of how linear algebra operations are integrated within PyTorch workflows and their impact on building effective machine learning models.